

Variable Neighbourhood Search (VNS)

Key Idea: systematically change neighbourhoods during search

Motivation:

- ▶ *recall*: changing neighbourhoods can help escape local optima
- ▶ a global optimum is locally optimal w.r.t. *all* neighbourhood structures
- ▶ *principle of VNS*: change the neighbourhood during the search

- ▶ main VNS variants
 - ▶ variable neighbourhood descent (VND, already discussed)
 - ▶ basic variable neighborhood search
 - ▶ reduced variable neighborhood search
 - ▶ variable neighborhood decomposition search

How to generate the various neighborhood structures?

- ▶ for many problems different neighborhood structures (local searches) exist / are in use
- ▶ use k -exchange neighborhoods; these can be naturally extended
- ▶ many neighborhood structures are associated with distance measures: define neighbourhoods in dependence of the distances between solutions

basic VNS

- ▶ uses neighborhood structures $\mathcal{N}_k, k = 1, \dots, k_{max}$
- ▶ iterative improvement in \mathcal{N}_1
- ▶ other neighborhoods are explored only randomly
- ▶ exploration in other neighborhoods are perturbations in the ILS sense
- ▶ perturbation is systematically varied
- ▶ acceptance criterion $\text{Better}(s^*, s^{*'})$

Basic VNS — Procedural view

procedure *basic VNS*

```
 $s_0 \leftarrow \text{GenerateInitialSolution, choose } \{\mathcal{N}_k\}, k = 1, \dots, k_{\max}$ 
repeat
   $s' \leftarrow \text{RandomSolution}(\mathcal{N}_k(s^*))$ 
   $s^{*'} \leftarrow \text{LocalSearch}(s')$  % local search w.r.t.  $\mathcal{N}_1$ 
  if  $f(s^{*'}) < f(s^*)$  then
     $s^* \leftarrow s^{*'}$ 
     $k \leftarrow 1$ 
  else
     $k \leftarrow k + 1$ 
until termination condition
end
```

Basic VNS — variants

- ▶ order of the neighborhoods
 - ▶ forward VNS: start with $k = 1$ and increase k by one if no better solution is found; otherwise set $k \leftarrow 1$
 - ▶ backward VNS: start with $k = k_{\max}$ and decrease k by one if no better solution is found
 - ▶ extended version: parameters k_{\min} and k_{step} ; set $k \leftarrow k_{\min}$ and increase k by k_{step} if no better solution is found
- ▶ acceptance of worse solutions
 - ▶ Skewed VNS: accept if

$$f(s^{*'}) - \alpha d(s^*, s^{*'}) < f(s^*)$$

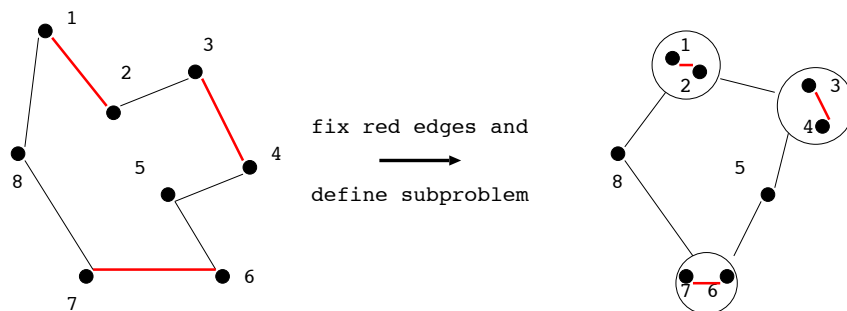
$d(s^*, s^{*'})$ measures distance between candidate solutions

Reduced VNS

- ▶ same as basic VNS except that no iterative improvement procedure is applied
- ▶ only explores randomly different neighborhoods
- ▶ goal: reach quickly good quality solutions for large instances

Variable Neighborhood Decomposition Search

- ▶ *central idea*
 - ▶ generate subproblems by keeping all but k solution components fixed
 - ▶ apply local search only to the k “free” components



- ▶ related approaches: POPMUSIC, MIMAUSA, etc.

VNDS — Procedural view

procedure *VNDS*

$s_0 \leftarrow \text{GenerateInitialSolution}$, choose $\{\mathcal{N}_k\}$, $k = 1, \dots, k_{\max}$

repeat

$s' \leftarrow \text{RandomSolution}(\mathcal{N}_k(s))$

$t \leftarrow \text{FixComponents}(s', s)$

$t^* \leftarrow \text{LocalSearch}(t)$ % local search w.r.t. \mathcal{N}_1

$s'' \leftarrow \text{InjectComponents}(t^*, s')$

if $f(s'') < f(s)$ **then**

$s \leftarrow s''$

$k \leftarrow 1$

else

$k \leftarrow k + 1$

until *termination condition*

end

relationship between ILS and VNS

- ▶ the two SLS methods are based on different underlying “philosophies”
- ▶ they are similar in many respects
- ▶ ILS appears to be in literature more flexible w.r.t. optimization of the interaction of modules
- ▶ VNS gives place to approaches like VND for obtaining more powerful local search approaches

Greedy Randomised Adaptive Search Procedures

Key Idea: Combine randomised constructive search with subsequent perturbative local search.

Motivation:

- ▶ Candidate solutions obtained from construction heuristics can often be substantially improved by perturbative local search.
- ▶ Perturbative local search methods typically often require substantially fewer steps to reach high-quality solutions when initialised using greedy constructive search rather than random picking.
- ▶ By iterating cycles of constructive + perturbative search, further performance improvements can be achieved.

Greedy Randomised “Adaptive” Search Procedure (GRASP):

While *termination criterion* is not satisfied:

- generate candidate solution s using
subsidiary greedy randomised constructive search
- perform *subsidiary local search* on s

Note:

Randomisation in *constructive search* ensures that a large number of good starting points for *subsidiary local search* is obtained.

Restricted candidate lists (RCLs)

- ▶ Each step of *constructive search* adds a solution component selected uniformly at random from a *restricted candidate list (RCL)*.
- ▶ RCLs are constructed in each step using a *heuristic function* h .
- ▶ RCLs based on *cardinality restriction* comprise the k best-ranked solution components. (k is a parameter of the algorithm.)
- ▶ RCLs based on *value restriction* comprise all solution components l for which $h(l) \leq h_{min} + \alpha \cdot (h_{max} - h_{min})$, where h_{min} = minimal value of h and h_{max} = maximal value of h for any l . (α is a parameter of the algorithm.)

Note:

- ▶ Constructive search in GRASP is 'adaptive':
Heuristic value of solution component to be added to given partial candidate solution r may depend on solution components present in r .
- ▶ Variants of GRASP without perturbative local search phase (aka *semi-greedy heuristics*) typically do not reach the performance of GRASP with perturbative local search.

Example: GRASP for SAT [Resende and Feo, 1996]

- ▶ **Given:** CNF formula F over variables x_1, \dots, x_n
- ▶ **Subsidiary constructive search:**
 - ▶ start from empty variable assignment
 - ▶ in each step, add one atomic assignment (i.e., assignment of a truth value to a currently unassigned variable)
 - ▶ heuristic function $h(i, v) :=$ number of clauses that become satisfied as a consequence of assigning $x_i := v$
 - ▶ RCLs based on cardinality restriction (contain fixed number k of atomic assignments with largest heuristic values)
- ▶ **Subsidiary local search:**
 - ▶ iterative best improvement using 1-flip neighbourhood
 - ▶ terminates when model has been found or given number of steps has been exceeded

GRASP has been applied to many combinatorial problems, including:

- ▶ SAT, MAX-SAT
- ▶ the Quadratic Assignment Problem
- ▶ various scheduling problems

Extensions and improvements of GRASP:

- ▶ reactive GRASP (e.g., dynamic adaptation of α during search)
- ▶ combinations of GRASP with Tabu Search and other SLS methods

Iterated Greedy

Key Idea: iterate over greedy construction heuristics through destruction and construction phases

Motivation:

- ▶ start solution construction from partial solutions to avoid reconstruction from scratch
 - ▶ keep features of the best solutions to improve solution quality
 - ▶ if few construction steps are to be executed, greedy heuristics are fast
- ▶ adding a subsidiary local search phase may further improve performance

Iterated Greedy (IG):

While *termination criterion* is not satisfied:

 | generate candidate solution s using
 | *subsidiary greedy constructive search*

While termination criterion is not satisfied:

 | $r := s$
 | apply *solution destruction* on s
 | perform *subsidiary greedy constructive search* on s

 | based on *acceptance criterion*,
 | keep s or revert to $s := r$

Note:

- ▶ subsidiary local search after solution reconstruction can substantially improve performance

Iterated Greedy (IG):

While *termination criterion* is not satisfied:

 generate candidate solution s using
 subsidiary greedy constructive search

While termination criterion is not satisfied:

$r := s$
 apply *solution destruction* on s
 perform *subsidiary greedy constructive search* on s
 perform *subsidiary local search* on s
 based on *acceptance criterion*,
 keep s or revert to $s := r$

Note:

- ▶ subsidiary local search after solution reconstruction can substantially improve performance

IG—main issues

- ▶ destruction phase
 - ▶ fixed vs. variable size of destruction
 - ▶ stochastic vs. deterministic destruction
 - ▶ uniform vs. biased destruction
- ▶ construction phase
 - ▶ not every construction heuristic is necessarily useful
 - ▶ typically, adaptive construction heuristics preferable
 - ▶ speed of the construction heuristic is an issue
- ▶ acceptance criterion
 - ▶ very much the same issue as in ILS

IG — enhancements

- ▶ usage of history information to bias destructive/constructive phase
- ▶ use lower bounds on the completion of a solution in the constructive phase
- ▶ combination with local search in the constructive phase
- ▶ use local search to improve full solutions
 \rightsquigarrow destruction / construction phases can be seen as a perturbation mechanism in ILS
- ▶ exploitation of constraint propagation techniques

Example: IG for SCP [Jacobs, Brusco, 1995]

- ▶ **Given:**
 - ▶ finite set $\mathbf{A} = \{a_1, \dots, a_m\}$ of objects
 - ▶ family $\mathbf{B} = \{B_1, \dots, B_n\}$ of subsets of \mathbf{A} that covers \mathbf{A}
 - ▶ weight function $w : \mathbf{B} \mapsto R^+$
- ▶ $\mathbf{C} \subseteq \mathbf{B}$ covers \mathbf{A} if every element in \mathbf{A} appears in at least one set in \mathbf{C} , i.e. if $\bigcup \mathbf{C} = \mathbf{A}$
- ▶ **Goal:**
 - ▶ find a subset $\mathbf{C}^* \subseteq \mathbf{B}$ of minimum total weight that covers \mathbf{A} .

Example: IG for SCP, continued ..

- ▶ assumption: all subsets from **B** are ordered according to nondecreasing costs
- ▶ construct initial solution using a greedy heuristic based on two steps
 - ▶ randomly select an uncovered object a_i
 - ▶ add the lowest cost subset that covers a_i
- ▶ the *destruction phase* removes a fixed number of $k_1|\mathbf{C}|$ subsets; k_1 is a parameter

- ▶ the *construction phase* proceeds as
 - ▶ build a candidate set containing subsets with cost of less than $k_2 \cdot f(\mathbf{C})$
 - ▶ compute cover value $\gamma_j = w_j/d_j$
 d_j : number of additional objects covered by adding subset b_j
 - ▶ add a subset with minimum cover value
- ▶ complete solution is post-processed by removing redundant subsets
- ▶ *acceptance criterion*: Metropolis condition from PII
- ▶ computational experience
 - ▶ good performance with this simple approach
 - ▶ more recent IG variants are state-of-the-art algorithms for SCP

- ▶ IG has been re-invented several times; names include
 - ▶ simulated annealing, ruin-and-recreate, iterative flattening, iterative construction search, large neighborhood search, ..
- ▶ close relationship to iterative improvement in large neighbourhoods
- ▶ analogous extension to greedy heuristics as ILS to local search
- ▶ for some applications so far excellent results
- ▶ can give lead to effective combinations of tree search and local search heuristics

Population-based SLS Methods

SLS methods discussed so far manipulate one candidate solution of given problem instance in each search step.

Straightforward extension: Use *population* (*i.e.*, set) of candidate solutions instead.

Note:

- ▶ The use of populations provides a generic way to achieve search diversification.
- ▶ Population-based SLS methods fit into the general definition from Chapter 1 by treating sets of candidate solutions as search positions.

Ant Colony Optimisation (1)

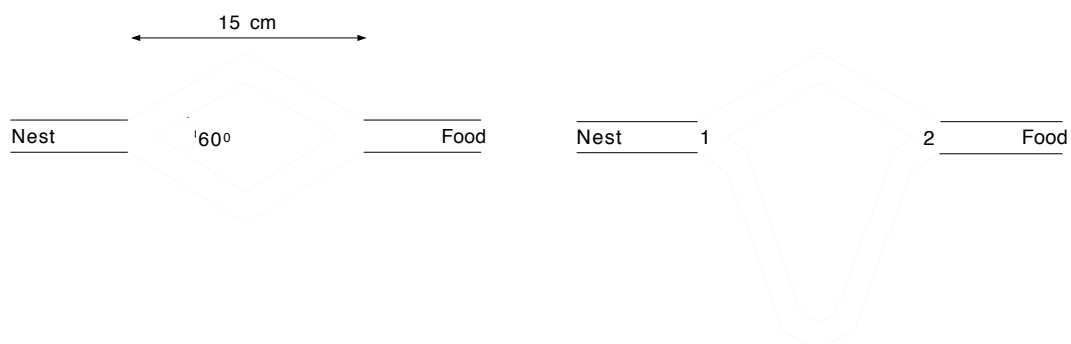
Key idea: Can be seen as population-based constructive approach where a population of agents – (*artificial*) *ants* – communicate via common memory – (*simulated*) *pheromone trails*.

Inspired by foraging behaviour of real ants:

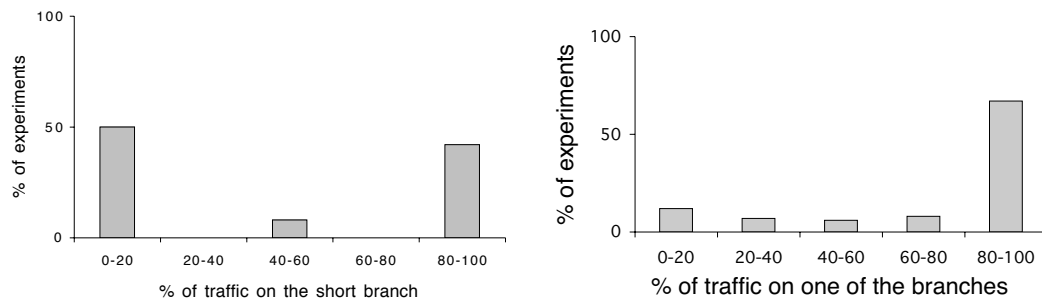
- ▶ Ants often communicate via chemicals known as *pheromones*, which are deposited on the ground in the form of trails. (This is a form of *stigmergy*: indirect communication via manipulation of a common environment.)
- ▶ Pheromone trails provide the basis for (stochastic) trail-following behaviour underlying, e.g., the collective ability to find shortest paths between a food source and the nest.

Double bridge experiments Deneubourg++

- ▶ laboratory colonies of *Iridomyrmex humilis*
- ▶ ants deposit pheromone while walking from food sources to nest *and* vice versa
- ▶ ants tend to choose, in probability, paths marked by strong pheromone concentrations



- ▶ equal length bridges: convergence to a single path
- ▶ different length paths: convergence to short path



- ▶ a stochastic model was derived from the experiments and verified in simulations
- ▶ functional form of transition probability

$$p_{i,a} = \frac{(k + \tau_{i,a})^\alpha}{(k + \tau_{i,a})^\alpha + (k + \tau_{i,a'})^\alpha}$$

- ▶ $p_{i,a}$: probability of choosing branch a when being at decision point i
 $\tau_{i,a}$: corresponding pheromone concentration
- ▶ good fit to experimental data with $\alpha = 2$

Towards artificial ants

- ▶ real ant colonies are solving *shortest path problems*
- ▶ ACO takes elements from real ant behavior to solve more complex problems than real ants
- ▶ In ACO, artificial ants are *stochastic solution construction procedures* that probabilistically build solutions exploiting
 - ▶ (artificial) *pheromone trails* that change at run time to reflect the agents' acquired search experience
 - ▶ *heuristic information* on the problem instance being solved

Ant Colony Optimisation (2)

Application to combinatorial problems:

[Dorigo et al. 1991, 1996]

- ▶ Ants iteratively construct candidate solutions.
- ▶ Solution construction is probabilistically biased by pheromone trail information, heuristic information and partial candidate solution of each ant.
- ▶ Pheromone trails are modified during the search process to reflect collective experience.

Ant Colony Optimisation (ACO):

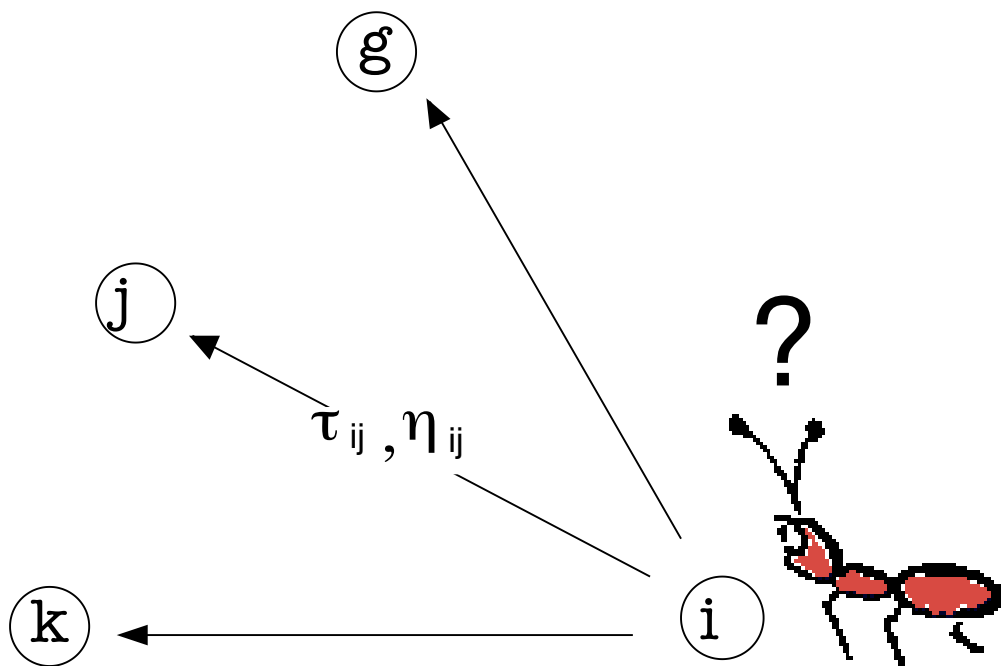
initialise pheromone trails

While termination criterion is not satisfied:

- generate population sp of candidate solutions
using *subsidiary randomised constructive search*
- perform *subsidiary local search* on sp
- update pheromone trails* based on sp

Note:

- ▶ In each cycle, each ant creates one candidate solution using a *constructive search procedure*.
- ▶ *Subsidiary local search* is applied to individual candidate solutions.
- ▶ All *pheromone trails* are initialised to the same value, τ_0 .
- ▶ *Pheromone update* typically comprises uniform decrease of all trail levels (*evaporation*) and increase of some trail levels based on candidate solutions obtained from construction + local search.
- ▶ *Termination criterion* can include conditions on make-up of current population, e.g., variation in solution quality or distance between individual candidate solutions.



Example: A simple ACO algorithm for the TSP (1)

(Variant of Ant System for the TSP [Dorigo *et al.*, 1991; 1996].)

- ▶ Search space and solution set as usual (all Hamiltonian cycles in given graph G).
- ▶ Associate pheromone trails τ_{ij} with each edge (i, j) in G .
- ▶ Use heuristic values $\eta_{ij} := 1/w((i, j))$.
- ▶ Initialise all weights to a small value τ_0 (parameter).
- ▶ *Constructive search*: Each ant starts with randomly chosen vertex and iteratively extends partial round trip ϕ by selecting vertex not contained in ϕ with probability

$$\frac{[\tau_{ij}]^\alpha \cdot [\eta_{ij}]^\beta}{\sum_{l \in N'(i)} [\tau_{il}]^\alpha \cdot [\eta_{il}]^\beta}$$

Example: A simple ACO algorithm for the TSP (2)

- ▶ *Subsidiary local search*: Perform iterative improvement based on standard 2-exchange neighbourhood on each candidate solution in population (until local minimum is reached).
- ▶ *Update pheromone trail levels* according to

$$\tau_{ij} := (1 - \rho) \cdot \tau_{ij} + \sum_{s' \in sp'} \Delta(i, j, s')$$

where $\Delta(i, j, s') := 1/f(s')$ if edge (i, j) is contained in the cycle represented by s' , and 0 otherwise.

Motivation: Edges belonging to highest-quality candidate solutions and/or that have been used by many ants should be preferably used in subsequent constructions.

Example: A simple ACO algorithm for the TSP (3)

- ▶ *Termination*: After fixed number of iterations (= construction + local search phases).

Note:

- ▶ Ants can be seen as walking along edges of given graph (using memory to ensure their tours correspond to Hamiltonian cycles) and depositing pheromone to reinforce edges of tours.
- ▶ Original Ant System did not include subsidiary local search procedure (leading to worse performance compared to the algorithm presented here)

<i>ACO algorithm</i>	<i>Authors</i>	<i>Year</i>	<i>TSP</i>
Ant System	Dorigo, Maniezzo, Colorni	1991	yes
Elitist AS	Dorigo	1992	yes
Ant-Q	Gambardella & Dorigo	1995	yes
<i>Ant Colony System</i>	Dorigo & Gambardella	1996	yes
<i>MMAS</i>	Stützle & Hoos	1996	yes
Rank-based AS	Bullnheimer, Hartl, Strauss	1997	yes
ANTS	Maniezzo	1998	no
Best-Worst AS	Cordón, et al.	2000	yes
Hyper-cube ACO	Blum, Roli, Dorigo	2001	no
Population-based ACO	Guntsch, Middendorf	2002	yes
Beam-ACO	Blum	2004	no

MAX-MIN Ant System

- ▶ extension of Ant System with stronger exploitation of best solutions and additional mechanism to avoid search stagnation
- ▶ *exploitation*: only the iteration-best or best-so-far ant deposit pheromone


$$\tau_{ij}(t+1) = (1 - \rho) \cdot \tau_{ij}(t) + \Delta\tau_{ij}^{best}$$

- ▶ frequently, a schedule for choosing between iteration-best and best-so-far update is used


- ▶ *stagnation avoidance*: additional limits on the feasible pheromone trails
 - ▶ for all $\tau_{ij}(t)$ we have: $\tau_{min} \leq \tau_{ij}(t) \leq \tau_{max}$
 - ▶ counteracts stagnation of search through aggressive pheromone update
 - ▶ heuristics for determining τ_{min} and τ_{max}
- ▶ *stagnation avoidance 2*: occasional pheromone trail re-initialization when \mathcal{MMAS} has converged
- ▶ *increase of exploration*: pheromone values are initialized to τ_{max} to have less pronounced differences in selection probabilities

solution quality versus time

d198




rat783




(typical parameter settings for high final solution quality)

behavior of ACO algorithms

average λ -branching factor



average distance among tours

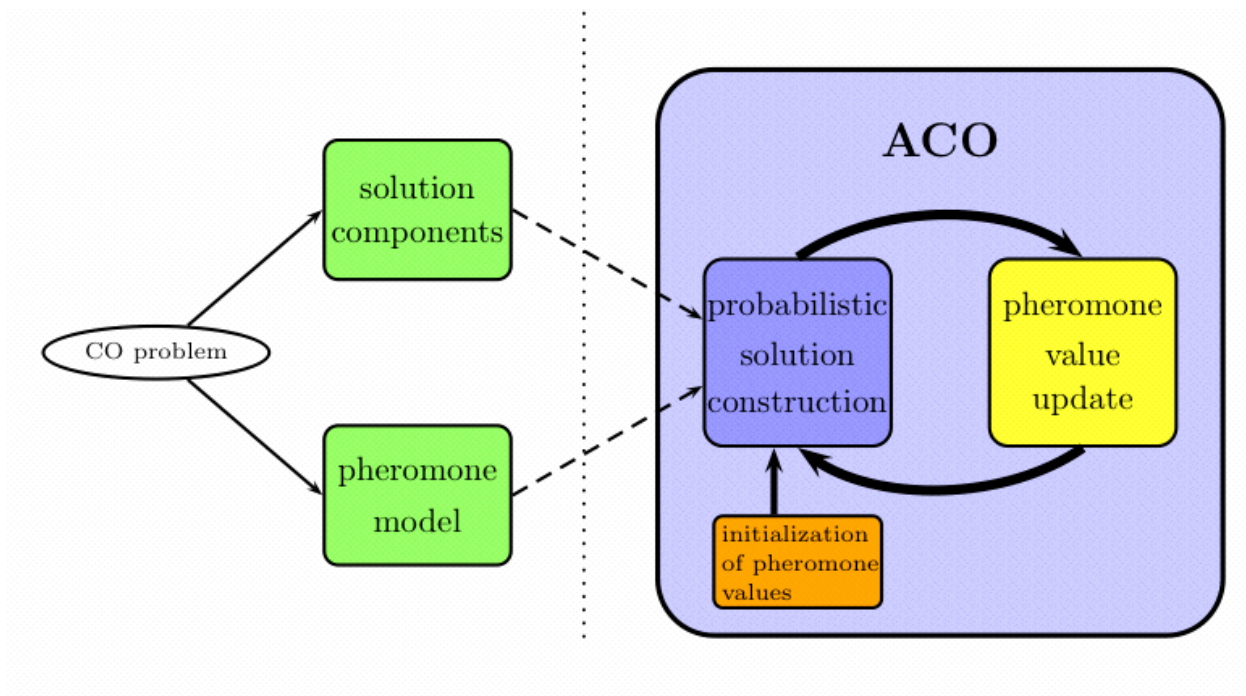


(typical parameter settings for high final solution quality)

Enhancements:

- ▶ use of look-ahead in construction phase;
- ▶ start of solution construction from partial solutions (memory-based schemes, ideas gleaned from iterated greedy);
- ▶ combination of ants with techniques from tree search such as
 - ▶ lower bounding information
 - ▶ combination with beam search
 - ▶ constraint programming (constraint propagation)

Applying ACO



copyright Christian Blum
Heuristic Optimization 2014

174

Ant Colony Optimisation ...

- ▶ has been applied very successfully to a wide range of combinatorial problems, including
 - ▶ the Open Shop Scheduling Problem,
 - ▶ the Sequential Ordering Problem, and
 - ▶ the Shortest Common Supersequence Problem;
- ▶ underlies new high-performance algorithms for *dynamic optimisation problems*, such as routing in telecommunications networks [Di Caro and Dorigo, 1998].

Note:

A general algorithmic framework for solving static and dynamic combinatorial problems using ACO techniques is provided by the *ACO metaheuristic* [Dorigo and Di Caro, 1999; Dorigo et al., 1999].

For further details on Ant Colony Optimisation, see the course on Swarm Intelligence or the book by Dorigo and Stützle [2004].

Evolutionary Algorithms

Key idea: Iteratively apply *genetic operators* *mutation*, *recombination*, *selection* to a population of candidate solutions.

Inspired by simple model of biological evolution:

- ▶ *Mutation* introduces random variation in the genetic material of individuals.
- ▶ *Recombination* of genetic material during reproduction produces *offspring* that combines features inherited from both *parents*.
- ▶ Differences in *evolutionary fitness* lead *selection* of genetic traits ('survival of the fittest').

Evolutionary Algorithm (EA):

determine initial population sp

While *termination criterion* is not satisfied:

generate set spr of new candidate solutions
by *recombination*

generate set spm of new candidate solutions
from spr and sp by *mutation*

select new population sp from
candidate solutions in sp , spr , and spm

Problem: Pure evolutionary algorithms often lack capability of sufficient *search intensification*.

Solution: Apply subsidiary local search after initialisation, mutation and recombination.

⇒ *Memetic Algorithms* (aka *Genetic Local Search*)

Memetic Algorithm (MA):

determine initial population sp

perform *subsidiary local search* on sp

While *termination criterion* is not satisfied:

 generate set spr of new candidate solutions
 by *recombination*

 perform *subsidiary local search* on spr

 generate set spm of new candidate solutions
 from spr and sp by *mutation*

 perform *subsidiary local search* on spm

select new population sp from
 candidate solutions in sp , spr , and spm

Initialisation

- ▶ *Often*: independent, uninformed random picking from given search space.
- ▶ *But*: can also use multiple runs of construction heuristic.


Recombination

- ▶ Typically repeatedly selects a set of *parents* from current population and generates *offspring* candidate solutions from these by means of *recombination operator*.
- ▶ *Recombination operators* are generally based on *linear representation* of candidate solutions and piece together *offspring* from fragments of *parents*.

Example: One-point binary crossover operator

Given two parent candidate solutions $x_1x_2 \dots x_n$ and $y_1y_2 \dots y_n$:

1. choose index i from set $\{2, \dots, n\}$ uniformly at random;
2. define offspring as $x_1 \dots x_{i-1}y_i \dots y_n$ and $y_1 \dots y_{i-1}x_i \dots x_n$.



Generalization: two-point, k -point, uniform crossover

Mutation

- ▶ *Goal:* Introduce relatively small perturbations in candidate solutions in current population + offspring obtained from *recombination*.
- ▶ Typically, perturbations are applied stochastically and independently to each candidate solution; amount of perturbation is controlled by *mutation rate*.
- ▶ Can also use *subsidiary selection function* to determine subset of candidate solutions to which mutation is applied.
- ▶ In the past, the role of mutation (as compared to recombination) in high-performance evolutionary algorithms has been often underestimated [Bäck, 1996].

Selection

- ▶ *Selection for variation*: determines which of the individual candidate solutions of the current population are chosen to undergo recombination and/or mutation
- ▶ *Selection for survival*: determines population for next cycle (*generation*) of the algorithm by selecting individual candidate solutions from current population + new candidate solutions obtained from *recombination*, *mutation* (+ *subsidiary local search*).
 - ▶ *Goal*: Obtain population of high-quality solutions while maintaining *population diversity*.
- ▶ Selection is based on evaluation function (*fitness*) of candidate solutions such that better candidate solutions have a higher chance of 'surviving' the selection process.

Selection (general)

- ▶ Many selection schemes involve probabilistic choices, using the idea that better candidate solutions have a higher probability of being chosen.
- ▶ examples
 - ▶ roulette wheel selection (probability of selecting a candidate solution s is proportional to its fitness value, $g(s)$)
 - ▶ tournament selection (choose best of k randomly sampled candidate solutions)
 - ▶ rank-based computation of selection probabilities
- ▶ the strength of the probabilistic bias determines the *selection pressure*

Selection (survival)

- ▶ *generational* replacement versus *overlapping populations*; (extreme case, *steady-state selection*)
- ▶ It is often beneficial to use *elitist selection strategies*, which ensure that the best candidate solutions are always selected.
- ▶ probabilistic versus deterministic replacement
- ▶ quasi-deterministic replacement strategies implemented by classical selections schemes from evolution strategies (a particular type of EAs)
 - ▶ λ : number offspring, μ : number parent candidate solutions
 - ▶ (μ, λ) strategy: choose best μ of $\lambda > \mu$ offspring
 - ▶ $(\mu + \lambda)$ strategy: choose best μ of $\mu + \lambda$ candidate solutions

Subsidiary local search

- ▶ Often useful and necessary for obtaining high-quality candidate solutions.
- ▶ Typically consists of selecting some or all individuals in the given population and applying an *iterative improvement procedure* to each element of this set independently.

Example: A memetic algorithm for SAT (1)

- ▶ *Search space*: set of all truth assignments for propositional variables in given CNF formula F ; *solution set*: models of F ; use *1-flip neighbourhood relation*; *evaluation function*: number of unsatisfied clauses in F .
- ▶ *Note*: truth assignments can be naturally represented as bit strings.
- ▶ Use population of k truth assignments; *initialise* by (independent) Uninformed Random Picking.

Example: A memetic algorithm for SAT (2)

- ▶ **Recombination**: Add offspring from $n/2$ (independent) one-point binary crossovers on pairs of randomly selected assignments from population to current population (n = number of variables in F).
- ▶ **Mutation**: Flip μ randomly chosen bits of each assignment in current population (*mutation rate* μ : parameter of the algorithm); this corresponds to μ steps of Uninformed Random Walk; mutated individuals are added to current population.
- ▶ **Selection**: Selects the k best assignments from current population (simple *elitist selection mechanism*).

Example: A memetic algorithm for SAT (3)

- ▶ **Subsidiary local search:** Applied after *initialisation*, *recombination* and *mutation*; performs *iterative best improvement* search on each individual assignment independently until local minimum is reached.
- ▶ **Termination:** upon finding model of F or after bound on number of cycles (*generations*) is reached.

Note: This algorithm does not reach state-of-the-art performance, but many variations are possible (few of which have been explored).

Problem representation and operators

- ▶ simplest choice of candidate solution representation: bitstrings
- ▶ advantage: application of simple recombination, mutation operators
- ▶ problems with this arise in case of
 - ▶ problem constraints (e.g. set covering, graph bi-partitioning)
 - ▶ “richer” problem representations are much better suited (e.g. TSP)
- ▶ *possible solutions*
 - ▶ application of representation- (and problem-) specific recombination and mutation operators
 - ▶ application of repair mechanisms to reestablish feasibility of candidate solutions

Memetic algorithm by Merz and Freisleben (MA-MF)

- ▶ one of the best studied MAs for the TSP
- ▶ first versions proposed in 1996 and further developed until 2001
- ▶ main characteristics
 - ▶ population initialisation by constructive search
 - ▶ exploits an effective LK implementation
 - ▶ specialised recombination operator
 - ▶ restart operators in case the search is deemed to stagnate
 - ▶ standard selection and mutation operators

MA-MF: population initialisation

- ▶ each individual of initial population constructed by a randomised variant of the greedy heuristic
 - Step 1: choose $n/4$ edges by the following two steps
 - ▶ select a vertex $v \in V$ uniformly at random among those that are not yet in partial tour
 - ▶ insert shortest (second-shortest) feasible edge incident to v with a probability of $2/3$ ($1/3$)
 - Step 2: complete tour using the greedy heuristic
- ▶ locally optimise initial tours by LK

MA-MF: recombination

- ▶ various specialised crossover operators examined (distance-preserving crossover, greedy crossover)
- ▶ crossover operators generate feasible offspring
- ▶ best performance by greedy crossover that borrows ideas from greedy heuristic
- ▶ one offspring is generated from two parents:
 1. copy fraction of p_e common edges to offspring
 2. add fraction of p_n new short edges not contained in any of the parents
 3. add fraction of p_c shortest edges from parents
 4. complete tour by greedy heuristic
- ▶ best performance for $p_e = 1$; $\mu/2$ pairs of tours are chosen uniformly at random from population for recombination

MA-MF: mutation, selection, restart, results

- ▶ *mutation* by (usual) double-bridge move
- ▶ *selection* done by usual $(\mu + \lambda)$ strategy
 - ▶ μ : population size, λ : number of new offspring generated
 - ▶ select μ lowest weight tours among $\mu + \lambda$ current tours for next iteration
 - ▶ take care that no duplicate tours occur in population
- ▶ *partial restart* by strong mutation
 - ▶ if average distance is below 10 or did not change for 30 iterations, apply random k -exchange move ($k = 0.1 \cdot n$) plus local search to all individuals except population-best one
- ▶ *results*: high solution quality reachable, though not fully competitive with state-of-the-art ILS algorithms for TSP

Memetic algorithm by Walters (MA-W)

- ▶ differs in many aspects from other MAs for the TSP
- ▶ main differences concern
 - ▶ solution representation by nearest neighbour indexing instead of permutation representation
 - ▶ usage of general-purpose recombination operators that may generate infeasible offspring
 - ▶ repair mechanism is used to restore valid tours from infeasible offspring
 - ▶ uses "only" a 3-opt algorithm as subsidiary local search

MA-W: solution representation, initialisation, mutation

- ▶ *solution representation* through nearest neighbour indexing
 - ▶ tour p represented as vector $s := (s_1, \dots, s_n)$ such that $s_i = k$ if, and only if, the successor of vertex u_i in p is k th nearest neighbour of u_i
 - ▶ leads, however, to some redundancies for symmetric TSPs
- ▶ *population initialisation* by choosing randomly nearest neighbour indices
 - ▶ three nearest neighbours selected with probability of 0.45, 0.25, 0.15, respectively
 - ▶ in remaining cases index between four and ten chosen uniformly at random
- ▶ *mutation* modifies nearest neighbour indices of randomly chosen vertices according to same probability distribution

MA-W: recombination, repair mechanism, results

- ▶ *recombination* is based on a slight variation of standard two-point crossover operator
- ▶ infeasible candidate solutions from crossover and mutation are repaired
- ▶ *repair mechanism* tries to preserve as many edges as possible and replaces an edge e by an edge e' such that $|w(e) - w(e')|$ is minimal
- ▶ *results* are interesting considering that "only" 3-opt was used as subsidiary local search; however, worse than state-of-the-art ILS algorithms

Tour merging

- ▶ can be seen as an extreme case of MAs
- ▶ exploits information collected by high-quality solutions from various ILS runs in a two phases approach
- ▶ *phase one*
 - ▶ generate a set T of very high quality tours for $G = (V, E, w)$
 - ▶ define subgraph $G' = (V, E', w')$, where E' contains all edges in at least one $t \in T$ and w' is original w restricted to E'
- ▶ *phase two*
 - ▶ determine optimal tour in G'
 - ▶ *Note:* general-purpose or specialised algorithm that exploit characteristics of T are applicable
- ▶ very high quality solutions can be obtained
 - ▶ optimal solution to TSPLIB instance d15112 in 22 days on a 500 MHz Alpha processor
 - ▶ new best-known solutions to instances brd14051 and d18512

Types of evolutionary algorithms (1)

- ▶ *Genetic Algorithms (GAs)* [Holland, 1975; Goldberg, 1989]:
 - ▶ have been applied to a very broad range of (mostly discrete) combinatorial problems;
 - ▶ often encode candidate solutions as bit strings of fixed length, which is now known to be disadvantageous for combinatorial problems such as the TSP.

Note: There are some interesting theoretical results for GAs (e.g., *Schema Theorem*), but – as for SA – their practical relevance is rather limited.

Types of evolutionary algorithms (2)

- ▶ *Evolution Strategies* [Rechenberg, 1973; Schwefel, 1981]:
 - ▶ originally developed for (continuous) numerical optimisation problems;
 - ▶ operate on more natural representations of candidate solutions;
 - ▶ use *self-adaptation* of perturbation strength achieved by *mutation*;
 - ▶ typically use *elitist deterministic selection*.
- ▶ *Evolutionary Programming* [Fogel *et al.*, 1966]:
 - ▶ similar to Evolution Strategies (developed independently), but typically does not make use of *recombination* and uses *stochastic selection* based on *tournament mechanisms*.